

Importing required libraries

```
In [115... import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from matplotlib import pyplot as plt
from matplotlib.colors import ListedColormap
import numpy as np
import pickle
import pandas as pd
import seaborn as sns
```

Reading data from dataset

```
In [63]: dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, :-1].values #all the numerical features
y = dataset.iloc[:, -1].values #label (or target)
```

```
In [64]: #dataset information
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID               400 non-null   int64
1   Gender                400 non-null   object
2   Age                   400 non-null   int64
3   EstimatedSalary       400 non-null   int64
4   Purchased             400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
In [65]: dataset.head()
```

```
Out[65]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [66]: dataset.tail()
```

```
Out[66]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

Data Preprocessing

```
In [67]: print(dataset.describe())
         #Missing Data Check
         print(dataset.isnull().sum())
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000
User ID	0			
Gender	0			
Age	0			
EstimatedSalary	0			
Purchased	0			

dtype: int64

This dataset does not have any missing data

```
In [68]: #Checking for duplicated values
         dataset.duplicated().sum()
```

```
Out[68]: 0
```

There are no duplicate values

```
In [69]: # Dropping Unecessary columns
         dataset.drop('User ID',axis=1,inplace=True)
         dataset.head()
```

Out[69]:

	Gender	Age	EstimatedSalary	Purchased
0	Male	19	19000	0
1	Male	35	20000	0
2	Female	26	43000	0
3	Female	27	57000	0
4	Male	19	76000	0

In [83]: *#Converting Categorical data to numeric data*
`dataset['Gender'].replace(to_replace=['Male', 'Female'], value=[1,2], inplace=True)`

Correlation between the features/variables are

In [84]: `cor = dataset.corr()
cor`

Out[84]:

	Gender	Age	EstimatedSalary	Purchased
Gender	1.000000	0.073741	0.060435	0.042469
Age	0.073741	1.000000	0.155238	0.622454
EstimatedSalary	0.060435	0.155238	1.000000	0.362083
Purchased	0.042469	0.622454	0.362083	1.000000

In [85]: `dataset['Purchased'].value_counts()`

Out[85]: Purchased
0 257
1 143
Name: count, dtype: int64

In [86]: `dataset.Age.describe()`

Out[86]: count 400.000000
mean 37.655000
std 10.482877
min 18.000000
25% 29.750000
50% 37.000000
75% 46.000000
max 60.000000
Name: Age, dtype: float64

Average age of the considered population is around 37

In [87]: `dataset.EstimatedSalary.describe()`


```
Out[87]: count      400.000000  
mean      69742.500000  
std       34096.960282  
min       15000.000000  
25%       43000.000000  
50%       70000.000000  
75%       88000.000000  
max       150000.000000  
Name: EstimatedSalary, dtype: float64
```

Average Salary of the population is around 70000

Checking if there is any relation between Age and Estimated salary

```
In [88]: plt.scatter(dataset.Age, dataset.EstimatedSalary)
```


```
Out[88]: <matplotlib.collections.PathCollection at 0x27e38479950>
```



We do not observe any major dependencies between Age and Salaries

```
In [89]: sns.catplot(x="Gender", col='Purchased', data=dataset, kind = 'count', palette='pastel')
```

```
Out[89]: <seaborn.axisgrid.FacetGrid at 0x27e36a4f950>
```



In [90]: `dataset.head()`

Out[90]:

	Gender	Age	EstimatedSalary	Purchased
0	1	19	19000	0
1	1	35	20000	0
2	2	26	43000	0
3	2	27	57000	0
4	1	19	76000	0

Training

In [136... *#Determining the columns that drive the decision of the "To purchase or not"*

```
feature_cols = ['Gender', 'Age', 'EstimatedSalary']
X = dataset[feature_cols]
Y = dataset['Purchased']
```

In [92]: *#splitting dataset into training and testing sets(75% training and 25% for testing)*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_
X_test
```

Out[92]:

	Gender	Age	EstimatedSalary
132	1	30	87000
309	2	38	50000
341	1	35	75000
196	2	30	79000
246	2	35	50000
...
146	2	27	96000
135	1	23	63000
390	1	48	33000
264	1	48	90000
364	1	42	104000

100 rows × 3 columns

In [93]: X_train

Out[93]:

	Gender	Age	EstimatedSalary
250	2	44	39000
63	1	32	120000
312	2	38	50000
159	2	32	135000
283	2	52	21000
...
323	2	48	30000
192	1	29	43000
117	1	36	52000
47	2	27	54000
172	2	26	118000

300 rows × 3 columns

In [94]: y_train

```
Out[94]: array([0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1,
 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0,
 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0,
 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1,
 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0,
 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0,
 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0,
 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1,
 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0,
 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0,
 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1,
 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0,
 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [95]: X_test
```

```
Out[95]:
```

	Gender	Age	EstimatedSalary
132	1	30	87000
309	2	38	50000
341	1	35	75000
196	2	30	79000
246	2	35	50000
...
146	2	27	96000
135	1	23	63000
390	1	48	33000
264	1	48	90000
364	1	42	104000

100 rows × 3 columns

Decision Tree Classifier

```
In [106... from sklearn import tree

tr = tree.DecisionTreeClassifier()
tr.fit(X_train, y_train)
y_pred = tr.predict(X_test)
y_pred
```

```
Out[106]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
                0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
                1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1,
                0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1,
                1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1], dtype=int64)
```

Feature Scaling is required to normalize the data from within a specified minimum and maximum range

This step is needed to normalize the range of independent variables/features from a minimum to a maximum range

Training the model and implementing the Test set

```
In [108... from sklearn import model_selection
kfold = model_selection.KFold(n_splits = 10)
tr = tree.DecisionTreeClassifier()
tr.fit(X_train, y_train)
results = model_selection.cross_val_score(tr, X_train, y_train, cv = kfold)
results
```

```
Out[108]: array([0.8      , 0.8      , 0.76666667, 0.76666667, 0.93333333,
                0.7      , 0.9      , 0.86666667, 0.93333333, 0.9      ])
```

```
In [109... tr_train_score = tr.score(X_train, y_train)

tr_test_score = tr.score(X_test, y_test)

print('Decision Tree Classifier Train Score is : ', tr_train_score)

print('Decision Tree Classifier Test Score is : ', tr_test_score)
```

Decision Tree Classifier Train Score is : 1.0

Decision Tree Classifier Test Score is : 0.91

Accuracy of the Classifier

```
In [110... from sklearn.metrics import accuracy_score

tr_acc = accuracy_score(y_test, y_pred)
print('accuracy_score : ', tr_acc)
```

accuracy_score : 0.91

```
In [111... from sklearn import metrics

tr_acc = 100*tr.score(X_test, y_test)
print('Decision Tree Classifier Predictions: \n', tr.predict(X_test), '\n Accuracy:
```


Decision Tree Classifier Predictions:

```
[0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 1 0 0 0 0 0 0 1 0 0 0 0
0 0 1 0 0 0 0 1 0 0 0 0 0 1 1 0 0 1 1 1 0 0 1 0 0 1 0 1 0 1 0 0 0 1 1 0 0 1
0 0 0 0 1 1 1 1 0 0 1 0 1 1 1 0 0 1 0 0 0 1 0 1 1 1]
```

Accuracy: 91.0 %


Confusion Matrix

```
In [141... from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[64  4]
 [ 3 29]]
```

```
In [151... from sklearn.metrics import confusion_matrix
cm_matrix = pd.DataFrame(data=cm, columns=['Actual_Approve', 'Actual_Deny'],
                        index=['Predict_Approve', 'Predict_Deny'])
sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

Out[151]: <Axes: >



Verification

```
In [124... cross_check = pd.DataFrame({'Actual' : y_test, 'Predicted' : y_pred})
cross_check
```

Out[124]:

	Actual	Predicted
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...
95	1	1
96	0	0
97	1	1
98	1	1
99	1	1

100 rows × 2 columns

In [132...

```

from sklearn.datasets import load_iris
from matplotlib import pyplot as plt
from sklearn import datasets, tree
from sklearn.tree import DecisionTreeClassifier
iris = load_iris()
#Prepare the data
x = iris.data
y = iris.target
clf = DecisionTreeClassifier(random_state=1234)
model = clf.fit(x, y)

```

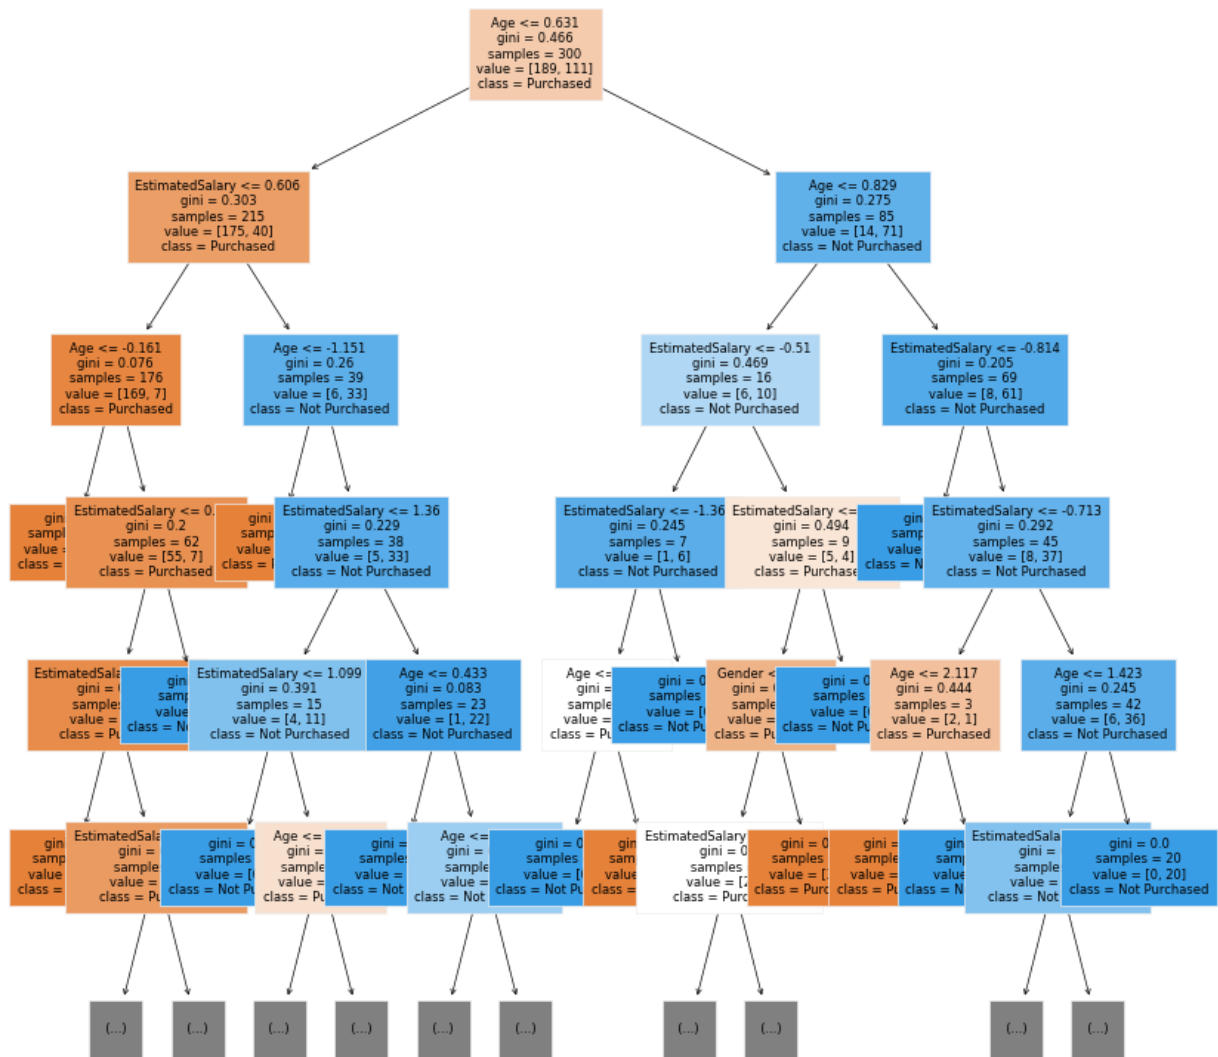
Plotting the decision tree

In [174...

```

fn=['Gender', 'Age', 'EstimatedSalary']
cn=['Purchased', 'Not Purchased']
fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (10,10), dpi=100)
tree.plot_tree(tr,
                feature_names = fn,
                class_names=cn,
                fontsize=6,
                max_depth = 5,
                filled = True);

```



KNN Model

```
In [137... from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train_scaled = sc.fit_transform(X_train)
X_test_scaled = sc.transform(X_test)
```

```
In [138... from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)
classifier.fit(X_train_scaled, y_train)
```

```
Out[138]: KNeighborsClassifier
KNeighborsClassifier()
```

```
In [139... y_pred = classifier.predict(X_test_scaled)
print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.reshape(len(y_test),1))
```

```
[[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[0 0]
[1 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[1 0]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[1 1]
[0 0]
[1 1]
[0 0]
[1 1]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 1]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[1 1]
[0 0]
[1 1]
[1 1]
[0 0]
[0 0]
[1 1]
[0 0]
[1 1]
[1 1]
[0 0]
[0 0]
[1 0]
[1 1]
[1 1]
```

```
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[1 1]
[0 0]
[1 1]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[1 1]
[1 1]
[1 0]
[0 0]
[0 0]
[1 1]
[0 1]
[0 0]
[1 1]
[1 1]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 1]
[0 0]
[1 1]
[1 1]
[1 1]]
```

```
In [140]: from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[64  4]
 [ 3 29]]
```

```
Out[140]: 0.93
```


Confusion Matrix

```
In [145... from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, y_pred))
```

```
[[64  4]
 [ 3 29]]
```

```
In [152... cm_matrix = pd.DataFrame(data=cm, columns=['Actual_Approve', 'Actual_Deny'],
                             index=['Predict_Approve', 'Predict_Deny'])
sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

Out[152]: <Axes: >




Plotting for K values ranging from 1 to 30

```
In [168... testAccuracy = []
trainAccuracy = []
for k in range(1,30):
    model=KNeighborsClassifier(n_neighbors=k)
    model.fit(X_train,y_train)
    trainAccuracy.append(model.score(X_train,y_train))
    testAccuracy.append(model.score(X_test,y_test))
```

```
In [166... from matplotlib import pyplot as plt, style
style.use('ggplot')
```

```
In [169... #create a plot using the information from the above loop
plt.figure(figsize=(12,6))
plt.plot(range(1,30),trainAccuracy,label="Train Score",marker="o",markerfacecolor="blue",markeredgecolor="blue",dash=[5,5])
plt.plot(range(1,30),testAccuracy,label="Test Score",marker="o",markerfacecolor="red",markeredgecolor="red",dash=[5,5])
plt.legend()
plt.xlabel("Number of Neighbors")
plt.ylabel("Score")
plt.title("Nbd Vs Score")
plt.show()
```



The optimal value hence obtained is 4

Creating a model using K = 4

```
In [171... from sklearn.pipeline import Pipeline
model_steps_20=[('sipStanderise',StandardScaler()),('shipModel',KNeighborsClassifie
pipelineModel=Pipeline(steps=model_steps_20)
pipelineModel.fit(X_train,y_train)
print("score is:"+ str(pipelineModel.score(X_train,y_train)))
print("*****")
pipelineModel.score(X_test,y_test)
predic_test_y=pipelineModel.predict(X_train)
print(pd.crosstab(y_train,predic_test_y))
```

```
score is:0.9133333333333333
*****
col_0    0    1
row_0
0       174   15
1        11  100
```

In []: